

# Day7

## Table of Contents

- [1. Scripts](#)
  - [1.1. compile script](#)
  - [1.2. run script](#)
- [2. MPI Type Struct with different blocklength](#)
  - [2.1. Compilation and Execution](#)
- [3. MPI Packing and Unpacking.](#)
  - [3.1. Functions](#)
  - [3.2. Syntax](#)
  - [3.3. Example Code](#)
  - [3.4. Explanation](#)
  - [3.5. Compilation and Execution](#)
- [4. MPI Pack Size, Probe, and Get Count](#)
  - [4.1. MPI Pack size](#)
    - [4.1.1. Syntax](#)
    - [4.1.2. Example](#)
  - [4.2. MPI Probe](#)
    - [4.2.1. Syntax](#)
  - [4.3. MPI Get count](#)
    - [4.3.1. Syntax](#)
  - [4.4. Example](#)
  - [4.5. Explanation](#)
  - [4.6. Compilation and Execution](#)
- [5. Task](#)
- [6. Task v2](#)

## 1. Scripts

### 1.1. compile script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cb
#spack load openmpi/c7kvqyq
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

inputFile=$1
outputFile="${1%.*}.out"      # extract the name of the file without extension and adding extension .out
#cmd=`mpicc $inputFile -o $outputFile`
cmd="mpicc $inputFile -o $outputFile"      # running code using MPI
echo "-----"
echo "Command executed: $cmd"
echo "-----"
$cmd

echo "Compilation successful. Check at $outputFile"
echo "-----"
```

## 1.2. run script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cb
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

cmd="mpirun -np $2 $1"
echo "-----"
echo "Command executed: $cmd"
echo "-----"
echo "#####"
echo "#####          OUTPUT          #####"
echo "#####"
echo
mpirun -np $2 $1
echo
echo "#####"
echo "#####          DONE          #####"
echo "#####"
```

## 2. MPI Type Struct with different blocklength

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <stddef.h>

typedef struct {
    int arr[3];
    double b;
    char c;
} my_struct;

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    my_struct data;
    MPI_Datatype struct_type;

    // Create the datatype for my_struct
    int blocklengths[3] = {2, 1, 1}; // Sending part of the array, the double, and the char
    MPI_Aint displacements[3];
    MPI_Datatype types[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};

    displacements[0] = offsetof(my_struct, arr);
    displacements[1] = offsetof(my_struct, b);
    displacements[2] = offsetof(my_struct, c);

    MPI_Type_create_struct(3, blocklengths, displacements, types, &struct_type);
    MPI_Type_commit(&struct_type);

    if (rank == 0) {
        data.arr[0] = 1;
        data.arr[1] = 2;
        data.arr[2] = 3;
        data.b = 3.14;
    }
}
```

```

    data.c = 'A';

    MPI_Send(&data, 1, struct_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 has struct: {arr = [%d, %d, %d], b = %.2f, c = %c}\n", data.arr[0], data.arr[1], data.arr[2], data.b, data.c);
} else if (rank == 1) {
    // Initialize the struct to zero
    data.arr[0] = data.arr[1] = data.arr[2] = 0;
    data.b = 0.0;
    data.c = '0';

    MPI_Recv(&data, 1, struct_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received struct: {arr = [%d, %d, %d], b = %.2f, c = %c}\n", data.arr[0], data.arr[1], data.arr[2], data.b, data.c);
}

MPI_Type_free(&struct_type);
MPI_Finalize();
return 0;
}

```

## 2.1. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_struct2.c
```

```

-----
Command executed: mpicc mpi_type_struct2.c -o mpi_type_struct2.out
-----
Compilation successful. Check at mpi_type_struct2.out
-----

```

- Run the program:

```
bash run.sh ./mpi_type_struct2.out 2
```

```

-----
Command executed: mpirun -np 2 ./mpi_type_struct2.out
-----

```

```
#####
#####          OUTPUT          #####
#####

Process 0 has struct: {arr = [1, 2, 3], b = 3.14, c = A}
Process 1 received struct: {arr = [1, 2, 0], b = 3.14, c = A}

#####
#####          DONE          #####
#####
```

This example demonstrates how to use `MPI\_Type\_create\_struct` to communicate complex data structures in MPI, specifically how to send parts of an array along with other fields.

## 3. MPI Packing and Unpacking

MPI provides mechanisms for packing and unpacking data into a contiguous buffer. This is useful for sending complex data structures without creating a custom MPI datatype. Instead, you manually pack the data into a buffer and then send the buffer.

### 3.1. Functions

- `MPI\_Pack`: Packs data of different types into a contiguous buffer.
- `MPI\_Unpack`: Unpacks data from a contiguous buffer.

### 3.2. Syntax

```
int MPI_Pack(const void *inbuf, int incount, MPI_Datatype datatype, void *outbuf, int outsize, int *position, MPI_Comm comm);
```

- `inbuf`: Input buffer containing data to be packed.
- `incount`: Number of elements in the input buffer.
- `datatype`: Datatype of each element in the input buffer.
- `outbuf`: Output buffer to contain packed data.
- `outsize`: Size of the output buffer.
- `position`: Current position in the output buffer (updated by MPI).

- ``comm``: Communicator.

```
int MPI_Unpack(const void *inbuf, int insize, int *position, void *outbuf, int outcount, MPI_Datatype datatype, MPI_Comm comm);
```

### 3.3. Example Code

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int a;
    double b;
    char c;
} my_struct;

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    my_struct data;
    int buffer_size, position;
    void *buffer;

    if (rank == 0) {
        data.a = 42;
        data.b = 3.14;
        data.c = 'A';

        // Calculate the buffer size required for packing
        MPI_Pack_size(1, MPI_INT, MPI_COMM_WORLD, &buffer_size);
        buffer_size += sizeof(double) + sizeof(char); // Adding the sizes of the other data types
        buffer = malloc(buffer_size);

        position = 0;
```

```

    MPI_Pack(&data.a, 1, MPI_INT, buffer, buffer_size, &position, MPI_COMM_WORLD);
    MPI_Pack(&data.b, 1, MPI_DOUBLE, buffer, buffer_size, &position, MPI_COMM_WORLD);
    MPI_Pack(&data.c, 1, MPI_CHAR, buffer, buffer_size, &position, MPI_COMM_WORLD);

    MPI_Send(buffer, position, MPI_PACKED, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent packed data\n");

    free(buffer);
} else if (rank == 1) {
    MPI_Status status;
    MPI_Probe(0, 0, MPI_COMM_WORLD, &status);
    MPI_Get_count(&status, MPI_PACKED, &buffer_size);

    buffer = malloc(buffer_size);
    MPI_Recv(buffer, buffer_size, MPI_PACKED, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    position = 0;
    MPI_Unpack(buffer, buffer_size, &position, &data.a, 1, MPI_INT, MPI_COMM_WORLD);
    MPI_Unpack(buffer, buffer_size, &position, &data.b, 1, MPI_DOUBLE, MPI_COMM_WORLD);
    MPI_Unpack(buffer, buffer_size, &position, &data.c, 1, MPI_CHAR, MPI_COMM_WORLD);

    printf("Process 1 received unpacked data: {a = %d, b = %.2f, c = %c}\n", data.a, data.b, data.c);
    free(buffer);
}

MPI_Finalize();
return 0;
}

```

### 3.4. Explanation

- **Initialization:** Initialize MPI, get the rank and size of the communicator.
- **Process 0:**
  - Initializes the `data` struct with values.
  - Calculates the buffer size required for packing the data using `MPI\_Pack\_size` and manually adds the sizes of the other data types.
  - Allocates memory for the buffer.
  - Packs each member of the struct into the buffer using `MPI\_Pack`.
  - Sends the packed buffer to process 1 using `MPI\_Send`.
  - Frees the buffer memory.
- **Process 1:**
  - Uses `MPI\_Probe` to get the size of the incoming message.

- Allocates memory for the buffer based on the received size.
- Receives the packed buffer from process 0 using `MPI\_Recv`.
- Unpacks each member of the struct from the buffer using `MPI\_Unpack`.
- Prints the unpacked data.
- Frees the buffer memory.
- **Finalize:** Finalize the MPI environment.

### 3.5. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_pack_unpack.c
```

```
-----
Command executed: mpicc mpi_pack_unpack.c -o mpi_pack_unpack.out
-----
Compilation successful. Check at mpi_pack_unpack.out
-----
```

- Run the program:

```
bash run.sh ./mpi_pack_unpack.out 2
```

```
-----
Command executed: mpirun -np 2 ./mpi_pack_unpack.out
-----
#####
#####          OUTPUT          #####
#####
#####

Process 0 sent packed data
Process 1 received unpacked data: {a = 42, b = 3.14, c = A}

#####
#####          DONE          #####
#####
```



This example demonstrates how to use ``MPI_Pack`` and ``MPI_Unpack`` to communicate complex data structures in MPI.

## 4. MPI Pack Size, Probe, and Get Count

### 4.1. MPI\_Pack\_size

`MPI_Pack_size` is used to calculate the size of the buffer needed to pack a message. This function helps ensure that the buffer you allocate is large enough to hold the packed data.

#### 4.1.1. Syntax

```
int MPI_Pack_size(int incount, MPI_Datatype datatype, MPI_Comm comm, int *size);
```

- ``incount``: Number of elements in the input buffer.
- ``datatype``: Datatype of each element in the input buffer.
- ``comm``: Communicator.
- ``size``: Pointer to the size of the packed message (output parameter).

#### 4.1.2. Example

Let's calculate the buffer size for packing an integer, a double, and a char.

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int size_int, size_double, size_char, total_size;
    MPI_Pack_size(1, MPI_INT, MPI_COMM_WORLD, &size_int);
    MPI_Pack_size(1, MPI_DOUBLE, MPI_COMM_WORLD, &size_double);
    MPI_Pack_size(1, MPI_CHAR, MPI_COMM_WORLD, &size_char);

    total_size = size_int + size_double + size_char;
    printf("Buffer size required for packing: %d bytes\n", total_size);
}
```

```

MPI_Finalize();
return 0;
}

```

- Compile the program:

```
bash compile.sh mpi_pack_size.c
```

```

-----
Command executed: mpicc mpi_pack_size.c -o mpi_pack_size.out
-----
Compilation successful. Check at mpi_pack_size.out
-----

```

- Run the program:

```
bash run.sh ./mpi_pack_size.out 2
```

```

-----
Command executed: mpirun -np 2 ./mpi_pack_size.out
-----
#####
#####          OUTPUT          #####
#####
#####
Buffer size required for packing: 13 bytes
Buffer size required for packing: 13 bytes

#####
#####          DONE          #####
#####

```

## 4.2. MPI\_Probe

MPI\_Probe allows you to probe for an incoming message without actually receiving it. This can be useful to determine the size of the message and allocate an appropriately sized buffer.

### 4.2.1. Syntax

```
int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status);
```

- ``source``: Rank of the source process (or ``MPI_ANY_SOURCE`` for any source).
- ``tag``: Message tag (or ``MPI_ANY_TAG`` for any tag).
- ``comm``: Communicator.
- ``status``: Status object that contains information about the message (output parameter).

## 4.3. MPI\_Get\_count

`MPI_Get_count` retrieves the number of elements of a specific datatype in a message. This function is often used after probing to determine the exact size of the received message.

### 4.3.1. Syntax

```
int MPI_Get_count(const MPI_Status *status, MPI_Datatype datatype, int *count);
```

- ``status``: Status object returned by ``MPI_Probe`` or ``MPI_Recv``.
- ``datatype``: Datatype of each element in the message.
- ``count``: Pointer to the number of received elements (output parameter).

## 4.4. Example

Let's combine ``MPI_Probe`` and ``MPI_Get_count`` to dynamically allocate a buffer for receiving a message.

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
```

```

int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

if (size < 2) {
    fprintf(stderr, "World size must be greater than 1 for this example\n");
    MPI_Abort(MPI_COMM_WORLD, 1);
}

if (rank == 0) {
    int data[5] = {1, 2, 3, 4, 5};
    MPI_Send(data, 5, MPI_INT, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data to process 1\n");
} else if (rank == 1) {
    MPI_Status status;
    MPI_Probe(0, 0, MPI_COMM_WORLD, &status);

    int count;
    MPI_Get_count(&status, MPI_INT, &count);

    int *buffer = (int*)malloc(count * sizeof(int));
    MPI_Recv(buffer, count, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    printf("Process 1 received %d integers:\n", count);
    for (int i = 0; i < count; i++) {
        printf("%d ", buffer[i]);
    }
    printf("\n");

    free(buffer);
}

MPI_Finalize();
return 0;
}

```

## 4.5. Explanation

- **MPI\_Pack\_size:**
  - This function is called three times to calculate the size required for packing an integer, a double, and a char.
  - The sizes are then summed to determine the total buffer size needed for packing.
- **MPI\_Probe:**

- Process 1 uses `MPI\_Probe` to check for an incoming message from process 0 without actually receiving it.
- The `status` object is filled with information about the message.
- **MPI\_Get\_count:**
  - `MPI\_Get\_count` is called to determine the number of integers in the received message using the `status` object from `MPI\_Probe`.
  - This allows process 1 to dynamically allocate a buffer of the appropriate size.

## 4.6. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_probe_get_count.c
```

```
-----
Command executed: mpicc mpi_probe_get_count.c -o mpi_probe_get_count.out
-----
Compilation successful. Check at mpi_probe_get_count.out
-----
```

- Run the program:

```
bash run.sh ./mpi_probe_get_count.out 2
```

```
-----
Command executed: mpirun -np 2 ./mpi_probe_get_count.out
-----
#####
#####          OUTPUT          #####
#####
#####

Process 0 sent data to process 1
Process 1 received 5 integers:
1 2 3 4 5

#####
#####          DONE          #####
#####
```

```
#####
```

## 5. Task

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <stddef.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    if (rank != 0) {
        MPI_Send(&rank, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
    } else {
        int data = 0;
        MPI_Recv(&data, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("I have received message from process %d\n", data);
    }

    MPI_Finalize();
    return 0;
}
```

```
bash compile.sh mpi_any_src_and_tag.c
```

```
-----
Command executed: mpicc mpi_any_src_and_tag.c -o mpi_any_src_and_tag.out
-----
```

```
Compilation successful. Check at mpi_any_src_and_tag.out
-----
```

```
bash run.sh ./mpi_any_src_and_tag.out 10
```

```
-----  
Command executed: mpirun -np 10 ./mpi_any_src_and_tag.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
#####  
  
I have received message from process 5  
  
#####  
#####          DONE          #####  
#####
```

## 6. Task v2

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#include <stdint.h>  
  
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);  
  
    int rank, size;  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);  
  
    if (size < 2) {  
        fprintf(stderr, "World size must be greater than 1 for this example\n");  
        MPI_Abort(MPI_COMM_WORLD, 1);  
    }  
  
    if (rank != 0) {  
        int data = 100;  
        MPI_Send(&data, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);  
    } else {  
        MPI_Status status;  
        int data = 0;  
        MPI_Recv(&data, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);  
    }  
}
```

```
    printf("I have received message from process %d\n", status.MPI_SOURCE);
}

MPI_Finalize();
return 0;
}
```

```
bash compile.sh mpi_any_src_and_tag.c
```

```
-----
Command executed: mpicc mpi_any_src_and_tag.c -o mpi_any_src_and_tag.out
-----
Compilation successful. Check at mpi_any_src_and_tag.out
-----
```

```
bash run.sh ./mpi_any_src_and_tag.out 10
```

```
-----
Command executed: mpirun -np 10 ./mpi_any_src_and_tag.out
-----
#####
#####                               OUTPUT                               #####
#####
#####

I have received message from process 8

#####
#####                               DONE                               #####
#####
#####
```

Author: Abhishek Raj  
Created: 2024-07-13 Sat 07:38