

# Day10

## Table of Contents

- [1. Scripts](#)
  - [1.1. compile script](#)
  - [1.2. run script](#)
- [2. Speedup, Amdahl's Law, and Performance Metrics in MPI](#)
  - [2.1. Performance Metrics](#)
  - [2.2. Speedup](#)
  - [2.3. Amdahl's Law](#)
  - [2.4. Measuring Performance in MPI](#)
- [3. Benchmarking in MPI](#)
  - [3.1. Benchmarking](#)
  - [3.2. Microbenchmarking](#)
  - [3.3. Macrobenchmarking](#)
  - [3.4. How Benchmarking is Done](#)
  - [3.5. Benchmarking Tools](#)
- [4. MPI Topology](#)
  - [4.1. Types of Topologies](#)
  - [4.2. Cartesian Topologies](#)
    - [4.2.1. Creating a Cartesian Topology](#)
  - [4.3. Graph Topologies](#)
    - [4.3.1. Creating a Graph Topology](#)
  - [4.4. Cartesian Topology Functions](#)
  - [4.5. Graph Topology Functions](#)
  - [4.6. Benefits of Using MPI Topologies](#)

## 1. Scripts

### 1.1. compile script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cb
#spack load openmpi/c7kvqyq
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

inputFile=$1
outputFile="${1%.*}.out"      # extract the name of the file without extension and adding extension .out
#cmd=`mpicc $inputFile -o $outputFile`
cmd="mpicc $inputFile -o $outputFile -lm -fopenmp"      # running code using MPI
echo "-----"
echo "Command executed: $cmd"
echo "-----"
$cmd

echo "Compilation successful. Check at $outputFile"
echo "-----"
```

## 1.2. run script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cbc
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

cmd="mpirun -np $2 $1"
echo "-----"
echo "Command executed: $cmd"
echo "-----"
echo "#####"
echo "#####          OUTPUT          #####"
echo "#####"
echo
mpirun -np $2 $1
echo
echo "#####"
echo "#####          DONE          #####"
echo "#####"
```

## 2. Speedup, Amdahl's Law, and Performance Metrics in MPI

### 2.1. Performance Metrics

When measuring the performance of MPI programs, several metrics are commonly used:

1. Execution Time:
  - The total time taken for the program to run.
2. Speedup:
  - As described above, it measures how much faster the parallel program runs compared to the serial version.
3. Efficiency:
  - Efficiency measures how effectively the processors are being utilized. It's defined as the speedup divided by the number of processors:

$$\text{Efficiency}(E) = \frac{S(p)}{p} = \frac{T_s}{p \times T_p}$$

4. Scalability:
  - Scalability refers to how well a parallel algorithm performs as the number of processors increases. It is evaluated using strong scaling and weak scaling:
    - Strong Scaling: Fixing the problem size and increasing the number of processors.
    - Weak Scaling: Increasing both the problem size and the number of processors proportionally.
5. Latency and Bandwidth:
  - Latency: The time taken to send a message from one process to another.
  - Bandwidth: The rate at which data can be transmitted.

### 2.2. Speedup

Speedup measures the improvement in performance of a parallel algorithm over its serial counterpart. It's calculated using the formula:

$$\text{Speedup}(S) = \frac{T_s}{T_p}$$

where:

- $T_s$  is the execution time of the serial algorithm.
- $T_p$  is the execution time of the parallel algorithm using  $p$  processors.

### 2.3. Amdahl's Law

Amdahl's Law provides a theoretical limit on the speedup that can be achieved by parallelizing a portion of an algorithm. It states that if a fraction  $f$  of the algorithm is inherently serial, then the maximum speedup  $S$  using  $p$  processors is:

$$S(p) = \frac{1}{f + \frac{1-f}{p}}$$

where:

- $f$  is the fraction of the algorithm that is serial (cannot be parallelized).
- $1-f$  is the fraction that can be parallelized.
- $p$  is the number of processors.

As  $p$  approaches infinity, the speedup approaches:

$$S_{\max} = \frac{1}{f}$$

This demonstrates that the speedup is limited by the serial portion of the algorithm.

### 2.4. Measuring Performance in MPI

To measure the performance of MPI programs, you can use the following methods:

## 1. Timing Functions:

- Use ``MPI_Wtime()`` to measure the wall-clock time before and after the parallel section of the code.

```
```c double start_time, end_time; start_time = MPI_Wtime();  
  
// Parallel code  
  
end_time = MPI_Wtime(); printf("Execution time: %f seconds\n", end_time - start_time);
```

2. Profiling Tools Use profiling tools like gprof, TAU, and VTune to analyze the performance of MPI programs.
3. MPI Profiling Interface (PMPI) MPI provides a profiling interface (PMPI) that allows you to intercept and measure MPI calls.
4. Performance Analysis Tools: Tools like Paraver, Vampir, HPCToolkit, and Scalasca can be used to visualize and analyze the performance of MPI programs.

## 3. Benchmarking in MPI

### 3.1. Benchmarking

Benchmarking is the process of measuring the performance of a system or application to evaluate its efficiency and effectiveness. In the context of MPI, benchmarking involves assessing the performance of MPI operations and parallel applications.

### 3.2. Microbenchmarking

Microbenchmarking focuses on measuring the performance of individual operations or small code segments. In MPI, microbenchmarking typically involves assessing the performance of basic MPI operations such as ``MPI_Send``, ``MPI_Recv``, ``MPI_Bcast``, etc.

### 3.3. Macrobenchmarking

Macrobenchmarking evaluates the performance of entire applications or larger code segments. It considers the overall performance and scalability of parallel applications, including computation, communication, and I/O.

### 3.4. How Benchmarking is Done

1. **Timing Functions:** Use functions like ``MPI_Wtime`` to measure execution time.
2. **Profiling Tools:** Utilize profiling tools to analyze performance.
3. **Performance Analysis Tools:** Use tools for in-depth performance analysis and visualization.

### 3.5. Benchmarking Tools

1. **OSU Micro-Benchmarks (OMB):** A suite of benchmarks for measuring MPI performance, focusing on latency, bandwidth, and collective operations.
2. **Intel MPI Benchmarks (IMB):** A set of benchmarks for evaluating the performance of MPI operations.
3. **HPC Challenge (HPCC):** A benchmark suite that measures the performance of HPC systems.
4. **SPEC MPI:** A benchmark suite designed to evaluate the performance of MPI-parallel, floating point, compute-intensive applications.

## 4. MPI Topology

MPI topologies provide a way to organize the processes in a communicator in a logical structure, which can enhance the performance of parallel applications by optimizing communication patterns.

### 4.1. Types of Topologies

1. **Cartesian Topologies:** Processes are arranged in a grid-like structure.
2. **Graph Topologies:** Processes are arranged in an arbitrary graph structure.

### 4.2. Cartesian Topologies

Cartesian topologies are useful for problems that have a natural grid structure.

### 4.2.1. Creating a Cartesian Topology

The function `MPI_Cart_create` is used to create a Cartesian topology.

#### 1. Syntax

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims, int *periods, int reorder, MPI_Comm *comm_cart);
```

#### 2. Parameters

- `comm_old`: The original communicator.
- `ndims`: Number of dimensions of the Cartesian grid.
- `dims`: Array specifying the number of processes in each dimension.
- `periods`: Array specifying whether the grid is periodic (wrap-around connections) in each dimension (logical array).
- `reorder`: Ranking may be reordered (1) or not (0).
- `comm_cart`: New communicator with Cartesian topology.

#### 3. Example Code

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    int dims[2] = {0, 0};
    MPI_Dims_create(size, 2, dims); // Automatically create a balanced 2D grid
    int periods[2] = {0, 0}; // No wrap-around connections
    MPI_Comm cart_comm;
    MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 0, &cart_comm);

    int coords[2];
    MPI_Cart_coords(cart_comm, rank, 2, coords);
    printf("Rank %d has coordinates (%d, %d)\n", rank, coords[0], coords[1]);

    // Cleanup
    MPI_Comm_free(&cart_comm);
}
```

```
MPI_Finalize();  
return 0;  
}
```

```
bash compile.sh mpi_cart.c
```

```
-----  
Command executed: mpicc mpi_cart.c -o mpi_cart.out -lm -fopenmp  
-----  
Compilation successful. Check at mpi_cart.out  
-----
```

```
bash run.sh ./mpi_cart.out 2
```

```
-----  
Command executed: mpirun -np 2 ./mpi_cart.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
  
Rank 1 has coordinates (1, 0)  
Rank 0 has coordinates (0, 0)  
  
#####  
#####          DONE          #####  
#####
```

## 4.3. Graph Topologies

Graph topologies are useful for problems with irregular communication patterns.

### 4.3.1. Creating a Graph Topology

The function ``MPI_Graph_create`` is used to create a graph topology.



## 1. Syntax

```
int MPI_Graph_create(MPI_Comm comm_old, int nnodes, int *index, int *edges, int reorder, MPI_Comm *comm_graph);
```

## 2. Parameters

- ``comm_old``: The original communicator.
- ``nnodes``: Number of nodes in the graph.
- ``index``: Array of integers describing the graph structure.
- ``edges``: Array of integers describing the graph structure.
- ``reorder``: Ranking may be reordered (1) or not (0).
- ``comm_graph``: New communicator with graph topology.

## 4.4. Cartesian Topology Functions

- ``MPI_Cart_create``: Creates a Cartesian topology.
- ``MPI_Cart_coords``: Determines the coordinates of a process in the Cartesian topology.
- ``MPI_Cart_rank``: Determines the rank of a process given its coordinates.
- ``MPI_Cart_shift``: Determines the source and destination ranks for shifts in a Cartesian topology.

## 4.5. Graph Topology Functions

- ``MPI_Graph_create``: Creates a graph topology.
- ``MPI_Graph_neighbors_count``: Determines the number of neighbors of a process.
- ``MPI_Graph_neighbors``: Determines the neighbors of a process.
- ``MPI_Graphdims_get``: Retrieves the number of nodes and edges in the graph topology.
- ``MPI_Graph_get``: Retrieves the graph structure.

## 4.6. Benefits of Using MPI Topologies

1. **Optimized Communication**: By organizing processes in a logical structure, communication can be optimized for better performance.
2. **Simplified Programming**: Topologies simplify the management of process coordinates and neighbor relationships.
3. **Improved Scalability**: Topologies can help applications scale more efficiently by reducing communication overhead.

Author: Abhishek Raj

Created: 2024-07-13 Sat 08:12