



Optimising Serial Code

Advancing Science with
Pawsey



Australian Government



Curtin University

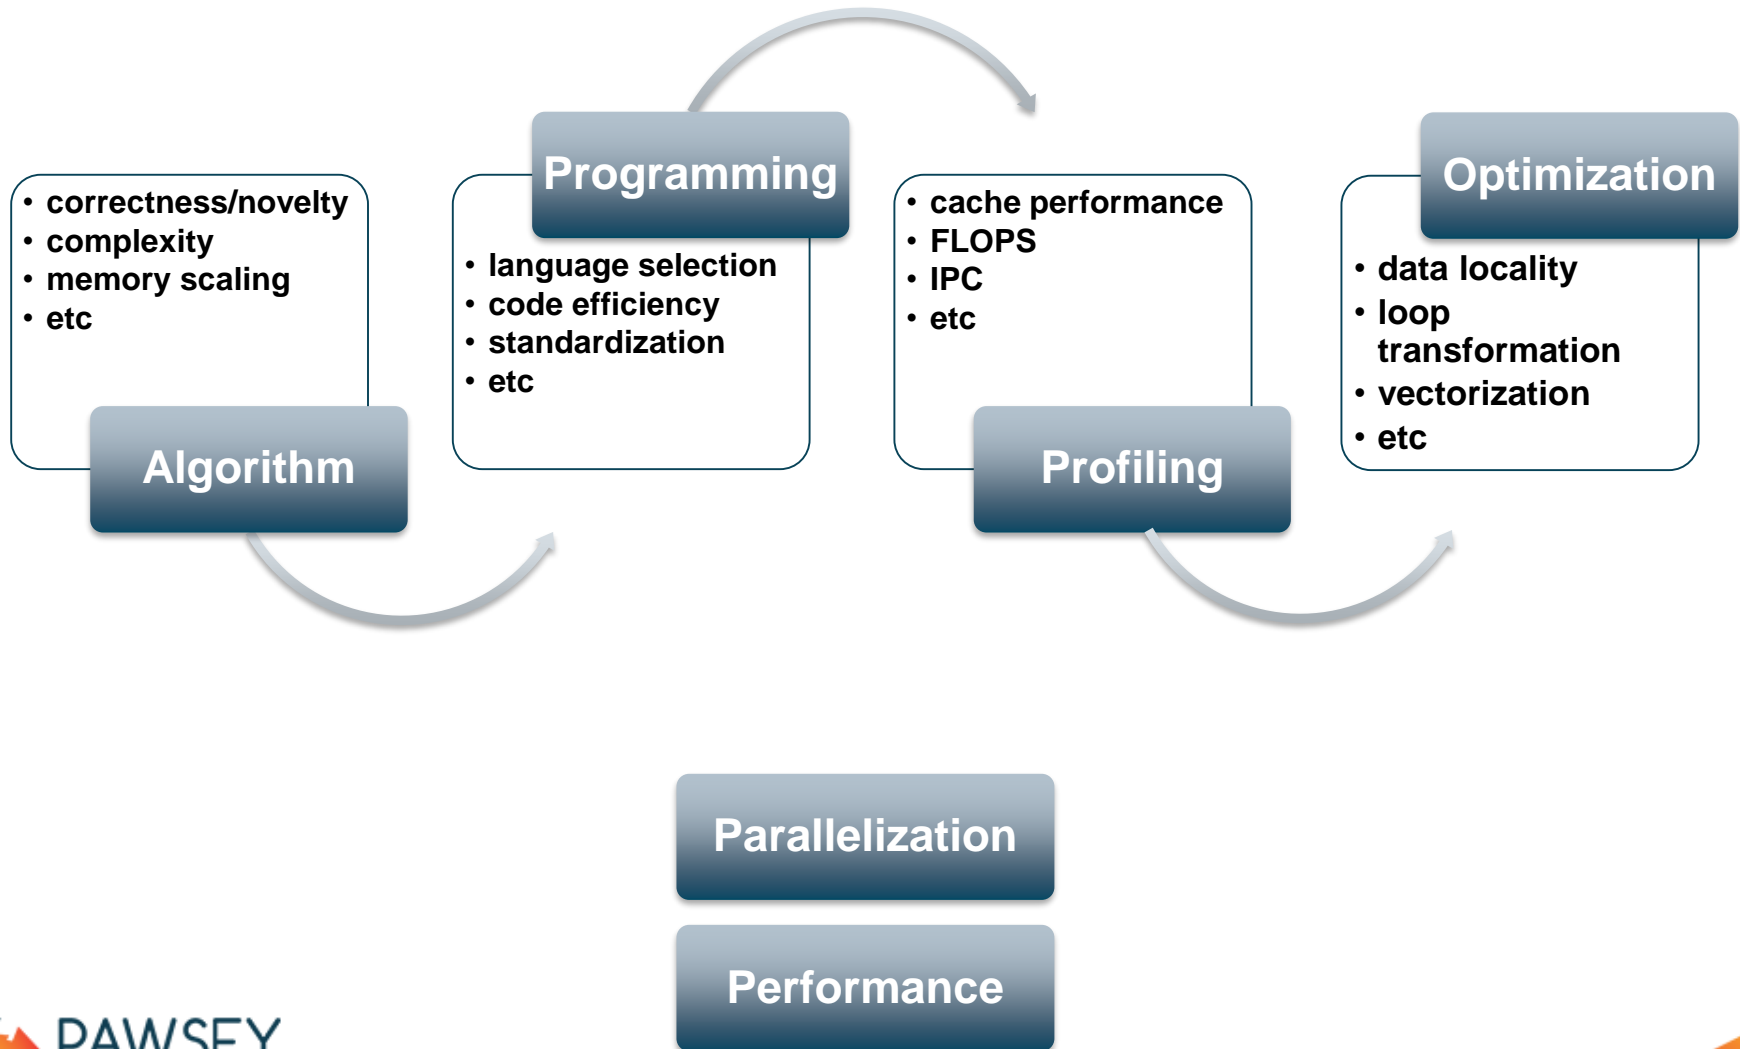


Learning Objectives

- Choose an algorithm for good performance
- Choose a language for good performance
- Understand the importance of standard conformance
- Write code that can be optimised for a modern CPU
- Locate bottlenecks in a serial code and address them



Course Roadmap



Motivation

- Science ***will*** be limited by the code and computer.
- Development effort may be secondary to code performance.
- Many gains come with small effort.
- Code optimisation ***can*** force good programming styles.

Is it worth the time?

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
HOW MUCH TIME YOU SHAVE OFF	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS

From XKCD webcomic: <https://xkcd.com/1205/>

Development Effort

1 Magnus day = 49 years on a dual-core desktop.

- ⇒ It is worth it to spend months of development time to save years/decades of runtime!
- ⇒ This might go against common opinions.

You still need to consider unit testing, ease of collaboration and code extension.



Know When to Stop Developing

Focus your effort!

1. Start with a good algorithm.
2. Write code with performance in mind.
3. Profile the code to determine where to spend your optimising effort.
4. Optimise the code, then profile again, and repeat until satisfactory performance achieved.

Know when to stop trying! Set an effort limit or a runtime goal.

MEASURING PERFORMANCE



Busy vs Effective

- Low processor utilisation **does** mean sub-optimal performance.
- High processor utilisation **does not** mean optimal performance.
- Why?
 - processor could be doing something inefficient
 - processor could be polling

Time

- What matters is real (your) time. Use the clock, we call this ***walltime***.
- You *can* use routines like **gettimeofday**, **MPI_Wtime** to measure program runtime.
- Don't use many timing calls to profile the code. This is slow and alters the profile. Instead use a profiler, which we will do later!
- CPU clock speed can vary between runs!

Timing routines

- C / C++

clock_t **clock**(void): Returns CPU time in “clock ticks” since initial call. (Use CLOCKS_PER_SEC to convert.). Resolution is microseconds.

time_t **time**(time_t *second): Returns real time in seconds since unix epoch 00:00:00 Jan 1, 1970. (Output argument “second” is also assigned with that value if it is not a NULL pointer.). Resolution is in seconds.

- Fortran

CPU_TIME(TIME): Returns the processor time in seconds since the start of the program through “TIME”, which is a REAL with INTENT(OUT). Resolution in milliseconds.

DATE_AND_TIME([date,time,zone,values]): Returns character and binary data on the real-time clock and date. Resolution is milliseconds.

- glibc

int **gettimeofday**(struct timeval *tp, struct timezone *tzp): Returns the current calendar time as the elapsed time since the epoch in the structure “tp”. Resolution is microseconds.

- MPI

double **MPI_Wtime**(void): Returns time in seconds since an arbitrary time in the past. Resolution is indicated by: double **MPI_Wtick**(void).

CHOOSING AN ALGORITHM



Algorithms

Many common algorithms have:

- a period of time
- increments within that period
- elements of interest
- ways to visit these elements
- a calculation <do something>

Algorithm A Grid	For t = t1 .. tn {over time increments} For i = x ₁ .. x _n For j = y ₁ .. y _n For k = z ₁ .. z _n <Do something>
Algorithm B N-body	For t = t1 .. tn {over time increments} For each element For each other element <Do something>

Complexity

At this stage, we can make a few observations:

Algorithm A will do $\#t.\#x.\#y.\#z$ calculations = $O(t^*n^3)$

Algorithm B will do $\#t.n.(n-1)$ calculations = $O(t^*n^2)$

... this is the **order** or algorithmic **complexity**.

Scaling

- Consider memory scaling as well as compute scaling. **Memory is finite.**

	Compute Time	Memory
3D-grid with time	$O(t \cdot n^3)$	$O(n^3)$
N-body	$O(t \cdot n^2)$	$O(n)$

- Other examples: QR tridiagonal eigensolver is $O(n^2)$ in memory, while MR³ tridiagonal eigensolver is $O(n)$ in memory.

Example Alternative Algorithm

- Instead of calculating every pairwise interaction (brute force) in the N-body problem, represent hierarchies of bodies with multipoles in a spacial decomposition (influence of very distant bodies is negligible).
- $O(N^2)$ interactions becomes $O(N \log N)$, even $O(N)$ for adaptive expansion.

N	Brute force $O(N^2)$	Multipole model $O(N \log N)$
1,000	1,000,000	7,000

Choosing an Algorithm

There may be multiple algorithms to solve a problem. Consider:

- E.g. $O(n^3)$ vs $O(n \log n)$ or Brute Force vs Clever approach with adequate results.
- Parallelisability constraints imposed by the algorithm.
- Domain decomposition often leads to good scaling and parallelisability. Load balancing may become an issue.

Prefactor

- If two algorithms have the same prefactor, does one have fewer expensive operations? E.g. square roots, exponentials, complex numbers, sin, cos etc.
- Algorithms that scale well might have a more expensive prefactor and be slower for small n . Perhaps use two algorithms and set a transition point based on n .

Rewrite your Maths

Minimise the number of array accesses and floating point operations, don't just go for elegant maths or copy a journal article verbatim. E.g.

- $\sum_{ij} X_{ij} Y_{ij}$ instead of $\text{tr}(X^T Y)$.
- $c \sum_i A_i$ instead of $\sum_i c A_i$.
- For symmetric matrices, you ***might*** gain from working with upper/lower triangles.
- Don't transpose a symmetric matrix.

Simple Tricks

- x^n – for small n make sure n is an integer. Then the compiler can change this to $x * x$.
- $\text{sqrt}(x)$ is generally faster than $x^{0.5}$.
- test $x^2 + y^2 < r^2$ instead of $\sqrt{x^2 + y^2} < r$.
- $\sum_{i \neq j} \frac{M_i M_j}{|r_{ij}|}$ if problem is symmetric then don't double the effort.
- Bring constants outside of loops, such as $\frac{1}{4 \pi \epsilon_0}$. The compiler *should* do this anyway.

CHOOSING A LANGUAGE



Choosing a language

Choose the right language for the job.

- Performance.
- Potential to use OpenMP, OpenACC, CUDA, and/or MPI.
- Flexible to new technologies.
- Portable.
- Available performance and debugging tools.
- Ease of programming.

Perhaps a hybrid approach, e.g. python/C or python/Fortran



Other Language Considerations

Our Goals:

- Want to convert algorithms into code (without too much effort).
- Want to help the compiler produce the fastest code (without too much effort).

Some Considerations:

- Array / matrix support
- Complex number support
- Aliasing
- Static typing



Aliasing

Aliasing is when some variables *might* refer to the same memory location. This introduces constraints on compiler optimisations. (e.g. prevents code reordering).

- Fortran

Assumes no aliasing. Avoid unnecessary use of pointers (stick to ***allocatable***).

- C

Use the ***restrict*** keyword, since C99 standard.

- C++

restrict etc are non-standard but in some compilers.
Not portable.

Avoid unnecessary use of pointers.

Static Typing

- With ***static typing***, the type is known at compile-time.
- Compiler can check that variables passed to functions are compatible with the functions.
- Permits optimisations at compile-time.
- Improves reproducibility, reduces errors.
- API is well documented for others.
- Fortran, C, C++ all support static typing.

Language performance

	Fortran	Julia	Python	R	Matlab	Octave	Mathematica	JavaScript	Go	SciLua	Java
	gcc 7.3.1	1.0.0	3.6.3	3.5.0	R2018a	4.2.2	11.3	V8 6.2.414	go1.9	1.0.0-b12	1.8.0_17
fibonacci	0.98	1.32	94.2	263	17.6	10045	132.1	3.51	1.80	1.18	3.63
parse_integer	6.87	2.19	19.76	50.37	178.2	574.4	22.66	5.03	0.96	0.97	3.17
quicksort	1.18	0.99	37.57	57.93	2.36	2221.3	44.48	4.28	1.24	1.56	2.98
mandelbrot	0.70	0.68	65.66	195.5	9.84	5812	18.29	1.134	0.77	1.00	1.42
pi_sum	1.00	1.01	14.77	11.69	1.01	317.5	1.45	1.08	1.00	0.99	1.08
matrix_statistics	1.54	1.63	17.73	20.97	8.09	46.24	7.49	13.97	6.08	1.70	5.02
matrix_multiply	1.15	0.97	1.18	8.26	1.16	1.21	1.18	31.77	1.43	1.08	8.07

Comparison from <https://julialang.org/benchmarks/>.

Times relative to C (gcc 7.3.1). Lower is better.

Ran on a single core (serial execution) on an Intel® Core™ i7-3960X 3.30GHz CPU with 64GB of 1600MHz DDR3 RAM, running openSUSE LEAP 15.0 Linux.

These are not optimal results, just how a typical researcher might program in those languages.

Fortran Compiler Comparison

	GCC	Intel	Cray
fibonacci	1	4.41	3.16
parse_integer	1	0.77	0.77
mandelbrot	1	0.50	-
quicksort	1	0.91	0.99
pi_sum	1	0.55	0.55
matrix_statistics	1	0.80	0.47
matrix_multiply	1	0.75	0.15

Comparison code perf.f90 from <https://julialang.org/benchmarks/>.

Runtimes normalised to the GNU runtimes. Lower is better.

Results on Magnus, all using “ftn -O3 perf.f90” under relevant compiler environments.

STANDARDS CONFORMANCE



Standard Conformance

- Your code will run on different systems and different compilers over the years.
- Standards conformance significantly improves the chance of **reproducibility**. (good for science!)
- Reproducibility means “apples vs apples” performance comparisons between systems and compilers, and **between profiling runs**.
- Do not rely on compiler extensions!

How to write portable code (1)

- Use your compiler to check.
 - `gcc -std=c99 -pedantic myfile.c`
 - `gcc -std=c++98 -pedantic myfile.cxx`
 - `gfortran -std=f95 -pedantic myfile.f90`
- Develop with multiple compilers. Cray compilers are very pedantic.
- Download a copy of the standard (or a draft).

How to write portable code (2)

- Try initialising variables to other values, do you get the same answer?.
 - `gfortran -finit-real=inf -finit-logical=false -finit-character=t ...`
- Try runtime checking. This is slower than compile-time checking.
- `gfortran -fcheck=all`

Portable Makefiles

- Don't hard code compiler-specific flags.
 - There is no standard for compiler flags, just for the code itself.
 - -Wall is not portable among compilers
- Use Makefile variables, make them easy to find and modify:

```
CFLAGS=-g -Wall
```

```
#CFLAGS=-O3
```

<http://www.gnu.org/software/make/manual/>

Exercise 1

Exercises are publicly available via Github:

git clone <https://github.com/PawseySC/Optimising-Serial-Code.git>

It's case sensitive. Download them!

Have a look to the various “Makefile”s.



Exercise 2: matmul (1)

In the exercise directory:

```
cd matrix && make matrix
```

This produces the executables `matrix.00`, `matrix.01`, `matrix.02`, `matrix.03` (Have a quick look at the Makefile).

Run them through the queue:

```
sbatch run_matrices.slurm
```

Output is in the SLURM output file `slurm-JOBID.out`

Exercise 2: matmul (2)

Compare the timings:


- What effect does optimisation level have on calls to the external math routine dgemm, to the intrinsic matmul, to manual looping of matrix multiplication?
- Is there a single best method for any matrix size at high optimisation?
- What is the main difference between matmul3 and matmul4.
- If you have time, try with a different compiler.

MODERN COMPUTERS AND OPTIMISATION



CPUs

- Most CPUs have multi-level caches, to reduce the time taken to access data.
- A CPU can do a lot of work in the time it takes to access main memory.



Data Locality

Location	Access time	Access time (cycles)
Register	<1ns	-
L1 cache	1ns	4
L2 cache	4ns	10
L3 cache	15-30ns	40-75
Memory	60ns	150
Solid state disk	50us	130,000
Hard disk	10ms	26,000,000
Tape	10sec	26,000,000,000

Source: https://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf

Cache: Access Patterns


- Sequential access results in a higher rate of cache hits



- Striding access has a low rate of hits, however modern processors can detect striding access patterns and pre-fetch cache lines



- Random access is typically the worst, with a low rate of hits and no ability to predict subsequent access locations




Translation Lookaside Buffer

- The Translation Page Table maps the memory seen by the program (Virtual Memory) into physical memory. It sits in main memory and is slow.
- The Translation Lookaside Buffer is a cache of recently used **mappings (not actual content)**. Like the main caches, aim to work within VM pages. These are typically 4kB.
- On a Cray you can use larger pages.
- Before **compiling**: `module load craype-hugepages2M` (size **does not** matter)
- **Runtime**: `module load craype-hugepagesXXXM` (size **does** matter)

Being cache-friendly

- Read and write to contiguous chunks of memory. Data is transferred in a *cache line*.
- Avoid *cache misses*.



Writing to Memory

- Don't store data in RAM if you don't need to. Use local temporary variables instead.
- These could be optimised into registers.
- In particular, don't use global variables as local temporary variables.
- Similarly, avoid using array sections for temporary storage.

Instruction Pipelining

- Modern CPUs can complete multiple Instructions Per Cycle (IPC), also known as Instructions Per Clock.
- An average for Xeon is 4. You need to keep the pipeline full to achieve this.

Loop Unrolling

Keep data in registers/cache via loop unrolling / blocking with inlining. This can also improve IPC.

e.g. $a(n) = \text{somefunc}(n) + a(n-1)$

Unroll with stride 2:

```
do n=1,nmax,2
```

```
    a(n)=somefunc(n) + a(n-1)
```

```
    a(n+1)=somefunc(n+1) + a(n)
```

```
end do
```

Loop Counts

- Inner loops should have high iteration counts, since loops themselves have a non-negligible cost.
- Counter to this, very small inner loops may get unrolled away. In this case do it manually.

Good

```
do j=1,10
  do n=1,10000
    work
  end do
end do
```

Bad

```
do n=1,10000
  do j=1,10
    work
  end do
end do
```

Knowing Loop Counts

- There is more potential for compiler optimisation when the loop count is known before the loop is started.
- Use “do”, “for” loops. Avoid “while” and “do ... while” loops.

Contiguous Memory

- Aim to work through contiguous chunks of memory. Avoid unnecessary striding.
- In **Fortran**, the consecutive elements of a column reside next to each other (**column-major order**).
- In **C/C++**, the consecutive elements of a row reside next to each other (**row-major order**)

Fortran

```
do j=1,10
  do i=1,10
    A(i,j)=something
```

C /C++

```
for (i=0;i<10;i++)
  for (j=0;j<10;j++)
    a[i][j]=something
```

Loop Blocking (1)

Loop blocking/tiling/strip-mining ...


This code potentially contains many cache misses.

```
do J=1,Jmax
  do I=1,Jmax
    A(I,J) = A(I,J) + B(J,I)
  end do
end do
```

A has stride 1, **B** has stride **Jmax**.

- Make the stride of **B** smaller so that small blocks of **A** and **B** sit in cache.
- If **I_{max}** and **J_{max}** are very large, neither A or B can sit in cache in whole.

Loop Blocking (2)



- The Cray Fortran compiler does this for you. Most other compilers currently do not.

Loop Blocking (3)

After applying loop blocking technique

```
do J=1,Jblksize,Jmax
  do I=1,Iblksize,Imax
    do JJ=J,J+Jblksize
      do II=I,I+Iblksize
        A(II,JJ) = A(II,JJ) + B(JJ,II)
      end do
    end do
  end do
end do
```

A has stride 1, **B** has stride **Jblksize**.

- Make the stride of **B** smaller so that small blocks of **A** and **B** sit in cache.
- Only **lblksize** x **Jblksize** of A or B sit in cache. Not exhausting cache.

Branching

- Remove branches from loops and change the loop bounds. Branches are bad for IPC and pre-emptive cache fetching.
- Avoid GOTO statements (in Fortran, C, C++). They *can* affect cache/register use.

Before:

```
do n=1,nmax
  if (n==1) a(n)=0
  a(n)=somefunc(n)
  if (n==nmax) a(n)=1
end do
```

After:

```
a(1)=0
a(nmax)=1
do n=2,nmax-1
  a(n)=somefunc(n)
end do
```

Loop Fission

- Too much work in loops means that registers and/or instruction cache may get exhausted.
- Perhaps only part of a loop is vectorisable (execute a number of loop-iterations at the same time).
- Break these loops up.

Before:

```
do n=1,nmax  
  lots of work  
  some I/O  
end do
```

After:

```
do n=1,nmax  
  lots of work  
end do  
do n=1,nmax  
  some I/O  
end do
```

Loop Fusion

Before:

```
do n=1,nmax
  a(n)=somefunc(n)
end do
do n=1,nmax
  b(n)=someotherfunc(n)
end do
```

After:

```
do n=1,nmax
  a(n)=somefunc(n)
  b(n)=someotherfunc(n)
end do
```

- This is usually used in conjunction with other techniques.
- Whether this is beneficial depends on the work inside the loops. Too much work may exhaust registers or cache.


Pointers

Avoid unnecessary use of pointers.

- Pointers ***might*** prevent some compiler optimisations. They should be fine if they have local scope.
- Pointers make it difficult to copy data to GPU memory.
- Hard to optimise cache use, e.g. with linked lists vs contiguous arrays.

MMX / 3DNOW / SSE / AVX

- These are Single-Instruction-Multiple-Data (SIMD) operations.
- These are not part of standard Fortran / C / C++. They may get deprecated over time.
- Compiler should insert these automatically. Write SIMD-friendly code.



Inlining

- Function calls take time. You can remove this time by placing the code into the calling routine.
- Compilers **can** inline code for you when using high optimization levels. This is not guaranteed, but you can make it more likely:
- In C / C++, put the code in the same file as the calling routine. Use static functions.
- In C99 / C++, use the **inline** keyword.
- In Fortran, put the code in the same file / module as the calling routine. Use the compiler directive *“forceinline”*.

COMPILER OUTPUT



Cray Compiler Output

Cray compiler will output annotated version of source file

```
ftn -rm mycode.f90
```

Outputs *mycode.lst*

Examine annotated file to figure out what's going on

%%% L o o p m a r k L e g e n d %%%

Primary Loop Type

Modifiers

A - Pattern matched

a - atomic memory operation

b - blocked

C - Collapsed

c - conditional and/or computed

D - Deleted

E - Cloned

F - Flat - No calls

f - fused

G - Accelerated

g - partitioned

I - Inlined

i - interchanged

M - Multithreaded

m - partitioned

n - non-blocking remote transfer

p - partial

R - Rerolling

r - unrolled

s - shortloop

V - Vectorized

w - unwound

Intel Compiler Output

- Optimisation reports.
- Compiler flag: `-qopt-report=3`
- Have a look at the man page for other values to `opt-report`.

Exercise 3: Cray compiler output

- Run:
`cd compiler_reports`
`make matrix.cray`
- Check out the manpage if needed
 - `man crayftn`
- Examine the output in `matrix.1st.0*`
- What has the compiler done with routine calls and loops?
- Can you identify the reason of the timing results from the previous exercise?

Exercise 4: Intel compiler output

- Run:

```
module swap PrgEnv-cray PrgEnv-intel  
cd compiler_reports  
make matrix.intel
```
- Seek help: `ifort -help reports`
- Examine the output in `matrix.optrpt.0*`
- Might be a bit too much information. Scale back the reporting options in `Makefile`

PROFILING



Profiling phases

- **Instrumentation:** compile the source code with extra compiler flags that enable the recording of performance-relevant events.
- **Measurement & analysis:** run the instrumented application on a **representative** test case. Usually the instrumented application is much slower than the original one.
- **Performance examination:** collect and analyse the measurement results.

Profilers

- On Cray supercomputers: Cray Tools

<https://pubs.cray.com/content/S-2376/7.0.0/cray-performance-measurement-and-analysis-tools-user-guide/craypat>

<https://support.pawsey.org.au/documentation/display/US/Profiling+with+Cray+Tools>

- Intel VTune: <https://software.intel.com/en-us/intel-vtune-amplifier-xe>
<https://support.pawsey.org.au/documentation/display/US/Profiling+with+Intel+VTune>

- Others: gprof, Arm MAP & Performance reports, etc.

<https://support.pawsey.org.au/documentation/display/US/User+Training>

<https://support.pawsey.org.au/documentation/display/US/Training+Material>

Profilers

- Profiling guide at Pawsey:

<https://support.pawsey.org.au/documentation/display/US/Profiling>

Profiling

Created by Maciej Cytowski on Sep 04, 2018

Pages in this section:

- [Profiling Introduction](#)
- [Basic profiling tools](#)
- [Profiling with gprof](#)
- [Profiling with Arm MAP](#)
- [Profiling with Arm Performance Reports](#)
- [Profiling with Cray Tools](#)
- [Profiling with Intel VTune](#)

Full Profiling with CrayPAT

Sampling experiment

Instrumentation

- `module load perftools`
- Compile code, using Cray compiler wrappers (`ftn`, `cc`, `CC`) & **preserving object (.o) files**
- `pat_build myapp`
 - Generates executable named `myapp+pat`

Measurement & analysis

- Run `./myapp+pat` as normal, this will generate an
- Output dir: `myapp+pat+XX+YYs/` (or `.xf` file for small runs)
- `pat_report myapp+pat+XX+YYs/ > myapp.sampling.report`
(this also generates `.ap2` file that can be viewed with Apprentice2, and a `build-options.apa` file to be used in a tracing experiment)

Performance examination

- Read `myapp.sampling.report` file
- or use Apprentice2 (with X11 forwarding activated: “`ssh -X`”):
`app2 myapp+pat+XX+YYs/ &`

Full Profiling with CrayPat

Tracing experiment

Instrumentation

- First, perform a sampling experiment to generate the file:
`myapp+pat+XX+YYs/build-options.apa`
- `pat_build -O myapp+pat+XX+YYs/build-options.apa`
 - Essentially `pat_build -w -T funcs -g grps -u myapp` (can be edited to change `-T funcs` and `-g grps`)
 - Generates executable named `myapp+apa`

Measurement & analysis

- Run `./myapp+apa` as normal, this will generate an
- Output dir: `myapp+apa+XX+YYt/` (or `.xf` file for small runs)
- `pat_report myapp+apa+XX+YYt/ > myapp.tracing.report`
(also generates `.ap2` file that can be viewed with Apprentice2)

Performance examination

- Read `myapp.tracing.report` file
- or use Apprentice2 (with X11 forwarding activated: “ssh -X”):
`app2 myapp+apa+XX+YYt/ &`

Exercise 5: profiling game of life (1)

Profile (sampling) the game_of_life code.

- `module load perftools`
- `cd game_of_life`
- `make game_of_life.cray`
- `pat_build game_of_life.03.cray`
- `sbatch run_game_profile.slurm`
- `pat_report game_of_life.03.cray+pat+XX-YYs/ > game_of_life.03.cray.sampling.report`

Exercise 5: profiling game of life (2)

Examine

`game_of_life.03.cray.sampling.report`

- Where is all the time spent? What occurs on these lines of the code?
- How good is our cache utilisation?
- Check optimisations in `game_of_life.03.lst`
- Try again with other levels of optimisation (edit the Makefile)

Exercise 6: profiling matrix multiplication (1) - **sampling**

Sampling experiment

- `module load perftools`
- `cd profiling`
- `ftn -c -O2 matrix.f90`
- `ftn -o matrix.O2 matrix.o`
- `pat_build matrix.O2`
- `sbatch run_matrix_profile.slurm`
- Generate the report:
`pat_report matrix.O2+pat+XX+YYs/ > matrix.O2.sampling.report`
- Examine `matrix.O2.sampling.report`
- You can also use `aprentice2: app2 matrix.O2+pat+XX+YYs/ &`

Exercise 6: profiling matrix (2) sampling

Table 2: Profile by Group, Function, and Line

Samp%	Samp	Imb. Samp	Imb. Samp%	Group Function Source Line
100.0%	566.0	--	--	Total
67.3%	381.0	--	--	BLAS
38.9%	220.0	--	--	gotoblas_dgemv_n_haswell
25.6%	145.0	--	--	__sci_dgemm_ bframe/crayblas/src/sci_gemm.c line.58
24.7%	140.0	--	--	USER
23.0%	130.0	--	--	test_matmul3\$timeit_ serialOptimisation/Optimising-Serial-Code/profiling/matrix.f90 line.190
22.3%	126.0	--	--	test_matmul1\$timeit_ serialOptimisation/Optimising-Serial-Code/profiling/matrix.f90
1.2%	7.0	--	--	
6.4%	36.0	--	--	ETC
4.6%	26.0	--	--	_RANF

Exercise 6: profiling matrix (3) sampling

Table 3: Program HW Performance Counter Data

Total

```
=====
Total
-----
Thread Time                                0.765879 secs
CPU_CLK_THREAD_UNHALTED:THREAD_P          1,710,445,839
CPU_CLK_THREAD_UNHALTED:REF_XCLK           50,726,596
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK       1,210,638
DTLB_STORE_MISSES:MISS_CAUSES_A_WALK       88,411
L1D:REPLACEMENT                           263,414,777
L2_RQSTS:ALL_DEMAND_DATA_RD                206,543,602
L2_RQSTS:DEMAND_DATA_RD_HIT                160,731,679
MEM_UOPS_RETIRED:ALL_LOADS                 1,948,873,169
CPU_CLK                                    3.37GHz
TLB utilization                            1,500.23 refs/miss      2.93 avg uses
D1 cache hit,miss ratios                   86.5% hits           13.5% misses
D1 cache utilization (misses)              7.40 refs/miss      0.92 avg hits
D2 cache hit,miss ratio                    82.6% hits           17.4% misses
D1+D2 cache hit,miss ratio                 97.6% hits           2.4% misses
D1+D2 cache utilization                    42.54 refs/miss      5.32 avg hits
D2 to D1 bandwidth                        16.074GiB/sec  13,218,790,528 bytes
=====
```


Exercise 6: profiling matrix (4) **sampling**

Table 6: Wall Clock Time, Memory High Water Mark

Process Time	Process HiMem (MBytes)	Total
0.801315	44.1	Total



Exercise 6: profiling matrix multiplication (5) - tracing

Tracing experiment

- Edit the build_options file to define tracing options:
`vim matrix.02+pat+XX-YYs/build_options.apa`

Have the following:

```
-g mpi,blas,io,heap  
-T test_matmul1$timeit_  
-T test_matmul2$timeit_  
-T test_matmul3$timeit_  
-T test_matmul4$timeit_  
-T test_matmul5$timeit_
```

- Build the executable with the defined tracing options:
`pat_build -O matrix.02+pat+XX-YYs/build_options.apa`

Exercise 6: profiling matrix (6) tracing

- Edit the job script: `vim run_matrix_profile.slurm`
Have: `expType=apa`
- `sbatch run_matrix_profile.slurm`
- Generate the report:
`pat_report matrix.02+apa+XX-YYt/ > matrix.02.tracing.report`
- Examine: `matrix.02.tracing.report`
- You can also use apprentice2: `app2 matrix.02+apa+XX+YYt/ &`
- Use the previous exercises of the matrix multiplication (timing and optimisation listing) to understand the results.
- Repeat the exercise with lower levels of optimisation.

Seek help

- `man pat_build`
- `man pat_report`
- `pat_help all . > all_pat_help`



Exercise 6: profiling matrix (7) tracing

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function
100.0%	0.691078	--	--	1,966.0	Total
75.5%	0.521487	--	--	6.0	USER
33.3%	0.230219	--	--	1.0	test_matmul4\$timeit_
19.4%	0.134307	--	--	1.0	test_matmul3\$timeit_
8.2%	0.056589	--	--	1.0	test_matmul1\$timeit_
7.3%	0.050596	--	--	1.0	test_matmul5\$timeit_
7.2%	0.049709	--	--	1.0	test_matmul2\$timeit_
21.6%	0.149372	--	--	1,302.0	BLAS
21.4%	0.147623	--	--	1.0	dgemm_
2.9%	0.020159	--	--	646.0	ETC
1.7%	0.011966	--	--	5.0	_END
1.1%	0.007688	--	--	384.0	query_cpu_mask

Exercise 6: profiling matrix (8) tracing

===== Observations and suggestions =====

D1 cache utilization:

1.7% of total execution time was spent in 1 functions with D1 cache hit ratios below the desirable minimum of 75.0%. Cache utilization might be improved by modifying the alignment or stride of references to data arrays in these functions.

D1 cache hit ratio	Time%	Function
-----------------------------	-------	----------

0.0%	1.7%	_END
------	------	------

D1 + D2 cache utilization:

All instrumented functions with significant execution time had combined D1 and D2 cache hit ratios above the desirable minimum of 80.0%.

TLB utilization:

1.1% of total execution time was spent in 1 functions with fewer than the desirable minimum of 200 data references per TLB miss. TLB utilization might be improved by modifying the alignment or stride of references to data arrays in these functions.

LS per TLB DM	Time%	Function
------------------	-------	----------

100.35	1.1%	query_cpu_mask
--------	------	----------------

Exercise 6: profiling matrix (9) tracing

```
=====
USER / test_matmul4$timeit_
-----
```

```
Time%                                33.3%
Time                                0.230219 secs
Imb. Time                           -- secs
Imb. Time%                           --
Calls                                4.344 /sec          1.0 calls
CPU_CLK_THREAD_UNHALTED:THREAD_P    785,898,184
CPU_CLK_THREAD_UNHALTED:REF_XCLK     22,709,160
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK 942,397
DTLB_STORE_MISSES:MISS_CAUSES_A_WALK 11,259
L1D:REPLACEMENT                     155,368,214
L2_RQSTS:ALL_DEMAND_DATA_RD          130,456,350
L2_RQSTS:DEMAND_DATA_RD_HIT          68,567,950
MEM_UOPS_RETIRED:ALL_LOADS           794,085,314
CPU_CLK                                3.46GHz
TLB utilization                       832.67 refs/miss          1.63 avg uses
D1 cache hit,miss ratios              80.4% hits              19.6% misses
D1 cache utilization (misses)         5.11 refs/miss          0.64 avg hits
D2 cache hit,miss ratio               60.2% hits              39.8% misses
D1+D2 cache hit,miss ratio            92.2% hits              7.8% misses
D1+D2 cache utilization               12.83 refs/miss          1.60 avg hits
D2 to D1 bandwidth                   33.776GiB/sec    8,349,206,400 bytes
Average Time per Call                  0.230219 secs
CrayPat Overhead : Time                0.0%
=====
```

Exercise 6: profiling matrix (10) tracing

```
=====
USER / test_matmul2$timeit_
-----
```

```
Time%                                7.2%
Time                                0.049709 secs
Imb. Time                           -- secs
Imb. Time%                           --
Calls                               20.117 /sec          1.0 calls
CPU_CLK_THREAD_UNHALTED:THREAD_P    173,201,511
CPU_CLK_THREAD_UNHALTED:REF_XCLK     5,211,090
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK  16,156
DTLB_STORE_MISSES:MISS_CAUSES_A_WALK  15,058
L1D:REPLACEMENT                     23,825,001
L2_RQSTS:ALL_DEMAND_DATA_RD          9,427,603
L2_RQSTS:DEMAND_DATA_RD_HIT         13,346,806
MEM_UOPS_RETIRE:ALL_LOADS            224,332,726
CPU_CLK                             3.32GHz
TLB utilization                      7,186.93 refs/miss    14.04 avg uses
D1 cache hit,miss ratios              89.4% hits          10.6% misses
D1 cache utilization (misses)         9.42 refs/miss       1.18 avg hits
D2 cache hit,miss ratio               100.0% hits          0.0% misses
D1+D2 cache hit,miss ratio            100.0% hits          0.0% misses
D1+D2 cache utilization               -- refs/miss         -- avg hits
D2 to D1 bandwidth                   11.304GiB/sec        603,366,592 bytes
Average Time per Call                  0.049709 secs
CrayPat Overhead : Time                0.0%
```

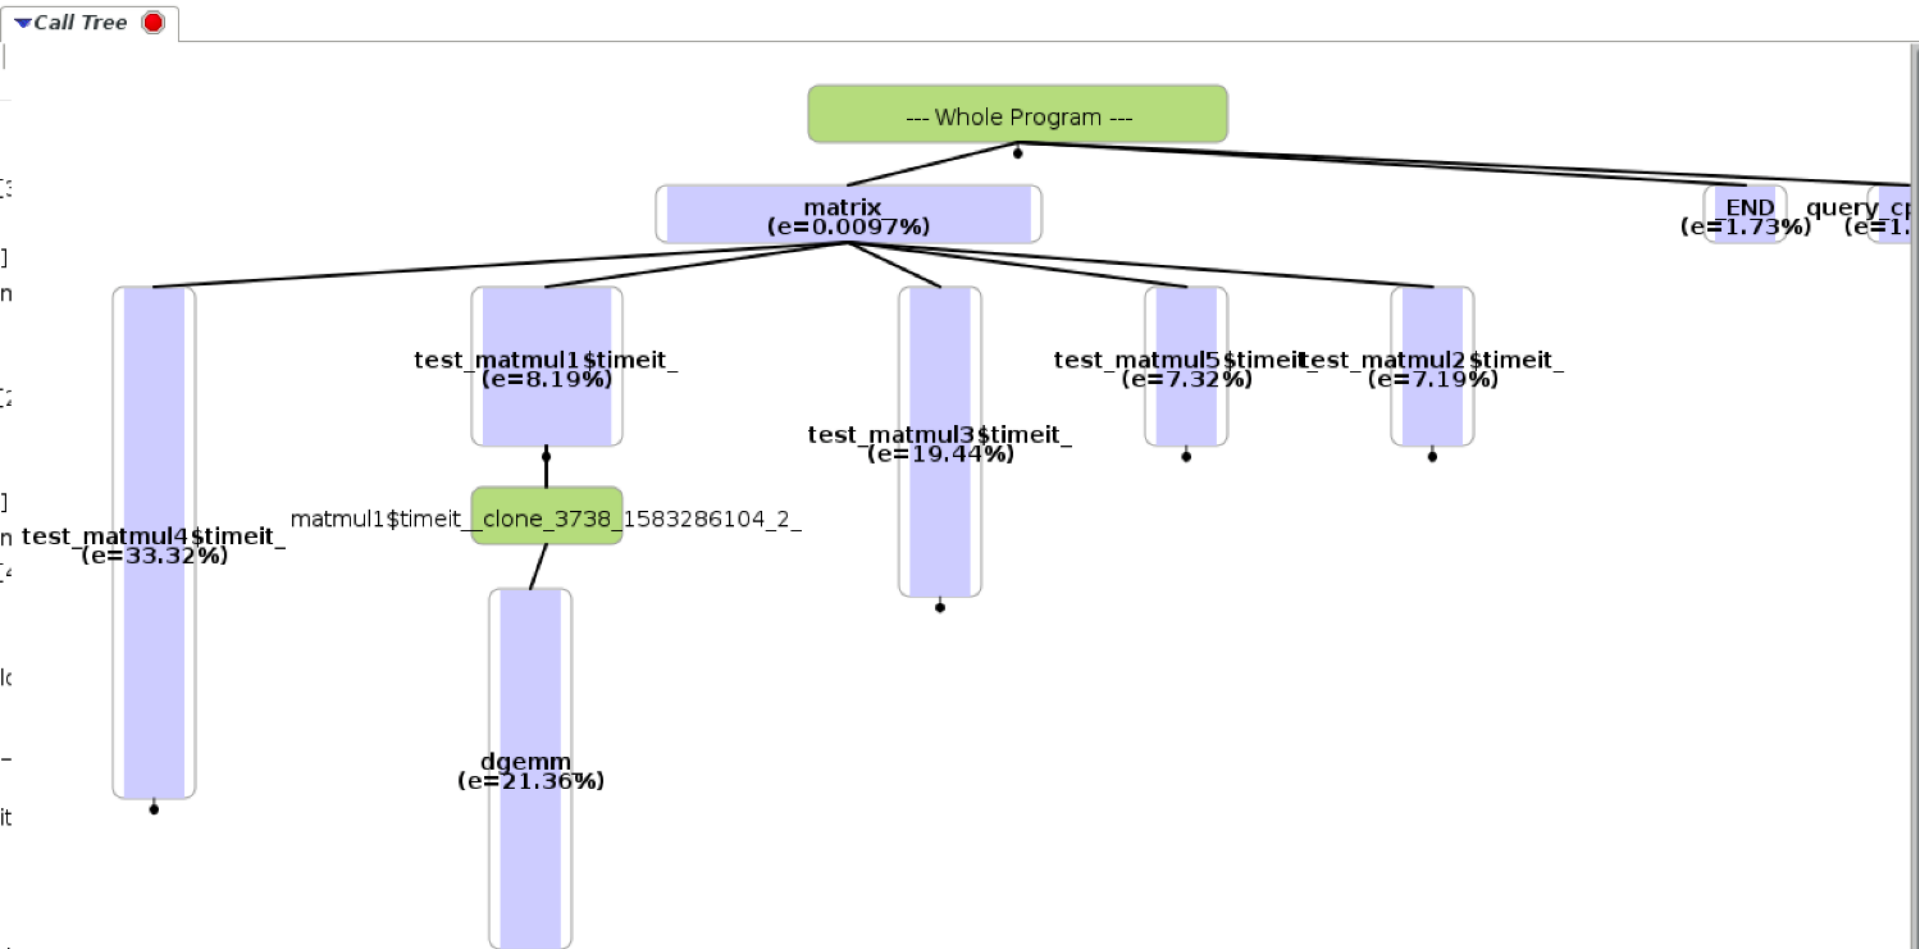
Exercise 6: profiling matrix (11) tracing

=====
BLAS / dgemm_

Time%		21.4%	
Time		0.147623	secs
Imb. Time		--	secs
Imb. Time%		--	
Calls	6.774 /sec	1.0	calls
CPU_CLK_THREAD_UNHALTED:THREAD_P		139,884,939	
CPU_CLK_THREAD_UNHALTED:REF_XCLK		4,202,951	
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK		26,457	
DTLB_STORE_MISSES:MISS_CAUSES_A_WALK		4,623	
L1D:REPLACEMENT		32,267,796	
L2_RQSTS:ALL_DEMAND_DATA_RD		21,348,184	
L2_RQSTS:DEMAND_DATA_RD_HIT		7,469,941	
MEM_UOPS_RETIRED:ALL_LOADS		193,382,571	
CPU_CLK	3.33GHz		
TLB utilization	6,222.09 refs/miss	12.15	avg uses
D1 cache hit,miss ratios	83.3% hits	16.7%	misses
D1 cache utilization (misses)	5.99 refs/miss	0.75	avg hits
D2 cache hit,miss ratio	57.0% hits	43.0%	misses
D1+D2 cache hit,miss ratio	92.8% hits	7.2%	misses
D1+D2 cache utilization	13.93 refs/miss	1.74	avg hits
D2 to D1 bandwidth	8.620GiB/sec	1,366,283,776	bytes
Average Time per Call		0.147623	secs
CrayPat Overhead : Time	0.0%		

Exercise 6: profiling matrix (12) tracing

From apprentice2: app2 matrix.02+apa+XX+YYt/ &



Exercise 6: profiling matrix (13) tracing

Table 3: Heap Stats during Main Program

Total

=====	
Total	

Tracked Heap HiWater MBytes	23.418
MBytes Not Tracked	0.000
Total Allocs	134
Allocs Not Tracked	0
Total Frees	128
Inferred Frees	0
Tracked Objects Not Freed	7
Tracked MBytes Not Freed	0.247
=====	



Exercise 6: profiling matrix multiplication (14) – tracing

Table 4: Heap Leaks during Main Program

Tracked MBytes Not Freed%	Tracked MBytes Not Freed	Tracked Objects Not Freed	Caller
100.0%	0.247	7	Total
28.8%	0.071	1	query_cpu_topology
28.1%	0.069	1	_GLOBAL__sub_I_eh_alloc.cc
25.6%	0.063	2	test_matmul1\$timeit_ matrix_
12.7%	0.031	1	resize_place_table
4.8%	0.012	1	_cray\$mt_init

Exercise 6: profiling matrix multiplication (15) – tracing

Table 5: File Input Stats by Filename

Read Time	Read MBytes	Read Rate MBytes/sec	Reads	Bytes/Call	File Name=!x/^/(proc sys)/
0.012854	0.051073	3.973411	1,393.0	38.45	Total

Table 6: File Output Stats by Filename (maximum 15 shown)

Write Time	Write MBytes	Write Rate MBytes/sec	Writes	Bytes/Call	File Name[max15]
0.000409	0.000796	1.947631	45.0	18.56	Total

0.000409	0.000796	1.947631	45.0	18.56	stdout
=====					

Exercise 6: profiling matrix multiplication (16) – **tracing**

Table 9: Wall Clock Time, Memory High Water Mark

Process Time	Process HiMem (MBytes)	Total
0.869646	46.1	Total



CODING HABITS



Global variables

- Avoid global variables unless necessary. They may make it difficult to convert to multithreaded code in further development.
- Pass variables through routine calls. (There is a slight performance overhead).
In Fortran arguments can be given intent(in), intent(out) attributes. This assists the compiler. Scoping in OpenMP becomes much easier. May assist in auto-threading by compilers.

Parentheses in Fortran

- Try to avoid parentheses in Fortran; they force an evaluation and prevent code arithmetic rearrangements. Use temporary variables instead.

- Compiler not permitted to rearrange this:

$a = 2 * (c + d) - 2 * e$

- Compiler allowed to rearrange this:

$tmp = c + d$

$a = 2 * tmp - 2 * e$

Fortran Array Notation

Fortran array notation is convenient and easy to read, but current compilers ***are likely to not optimise*** them well.

Some compilers are unlikely to fuse these operations:

```
A(:, :)=1.0  
C(:, :)=A(:, :)+B(:, :)
```

In the meantime:

```
do j,1,n  
  do i=1,m  
    A(i,j)=1.0  
    C(i,j)=A(i,j)+B(i,j)  
  end do  
end do
```

The Cray compiler does fuse array notation!

Low Level Object Oriented

Low level Object Oriented programming has the potential for poor performance. E.g. the below strided (not contiguous) memory accesses.

```
type atom_type
  integer :: atomic_number
  double precision :: mass
  double precision, dimension(3) :: position
end type atom_type

type(atom_type), dimension(n_atoms) :: atom_list

total_mass=0
do i=1,n_atoms
  total_mass=total_mass+atom_list(i)%mass
end do
```

Low Level Object Oriented (2)

Less organised but faster code:

```
integer, dimension(n_atoms) :: atomic_numbers  
double precision, dimension(n_atoms) :: masses  
double precision, dimension(3,n_atoms) :: positions  
  
total_mass=sum(masses)
```

Set a level for the trade-off between maintainability/extensibility and performance.

Special Case Code

Assume we have a code that handles arrays of varying length, and;

- the code creates temporary arrays;
- in practice an array of length 1 is the most common situation.

Optimisation: write a separate routine for arrays of length 1, and use temporary variables rather than temporary arrays.

I/O

- Often the processor is doing little while waiting for I/O.
- Ways to reduce I/O overhead:
 - Use buffering (or don't turn it off or flush).
 - Output in binary, not formatted text.
 - Use I/O libraries.
- Hierarchical Data Format (HDF5) is the name of a set of file formats and libraries designed to store and organize large amounts of numerical data.

Exercise 7: I/O

Observe the effects of I/O techniques on performance.

- `module load cray-hdf5`
- `cd iobench && make iobench_hdf5`
- `sbatch run_iobench_hdf5.slurm`

Look at the SLURM output file.

Use version control

- Some of your attempts at optimisation will need to be undone. E.g. due to:
 - Incorrect results.
 - Slower performance.
- Use version control software. E.g. git, subversion. You should be using this anyway.
- Use informative comments in check-in.

MATH LIBRARIES



Popular libraries

- BLAS: basic linear algebra such as matrix-vector or matrix-matrix operations.
- LAPACK:
 - Simple matrix/vector operations
 - Linear equations solvers
 - Linear least squares
 - Eigensolvers
 - Singular value decomposition
 - Real + Complex
- FFTW: fast fourier transforms, real/complex

Optimised vendor versions available. e.g. Intel MKL, Cray Libsci, SGI SCSL, IBM ESSL. Some are multi-threaded.

Other libraries

- LAPACK - better scaling eigensolver (MRRR algorithm)
- PARPACK – sparse eigensolver
- MUMPS – parallel sparse direct solver
- Hypre – parallel linear solver
- Scotch – graph partitioning
- SuperLU – parallel sparse linear solver
- available at cray-tpsl

Intel Math Libraries

Intel MKL includes BLAS, LAPACK and FFT libraries. It consists of multiple libraries – use the Intel advisor to work out the compiler link options:

http://software.intel.com/sites/products/mkl/MKL_Link_Line_Advisor.html

Example output:

```
-L$(MKLROOT)/lib/intel64 -lmkl_intel_lp64 -  
lmkl_sequential -lmkl_core -lpthread -lm
```

`$(MKLROOT)` is set by “module load intel”

PetSc / Slepc

- It's an attempt at a black box solver suite. (makes it easy to swap between various solvers in various libraries). It links to common libraries.
- C/C++ and Fortran interfaces
- Linear equation solvers
- Eigensolvers
- Dense and Sparse
- various finite element solver add-ons

Finish

- What's next?:
 - Come to the parallel course.
 - Read optimisation guides from vendors. Intel and Cray in particular.
- Slides are available at <https://support.pawsey.org.au/documentation/display/US/Training+Material>