


OPENMP OFFLOADING

GPU Work

What is offloading ?

- Offloading is the process of transferring a portion of a computation from a central processor (such as a CPU) to a specialized accelerator (such as a coprocessor , FPGA or GPU) or to a remote machine.
- It can be used to improve the performance of a computation



What is the Difference
Between GPU and
CPU?



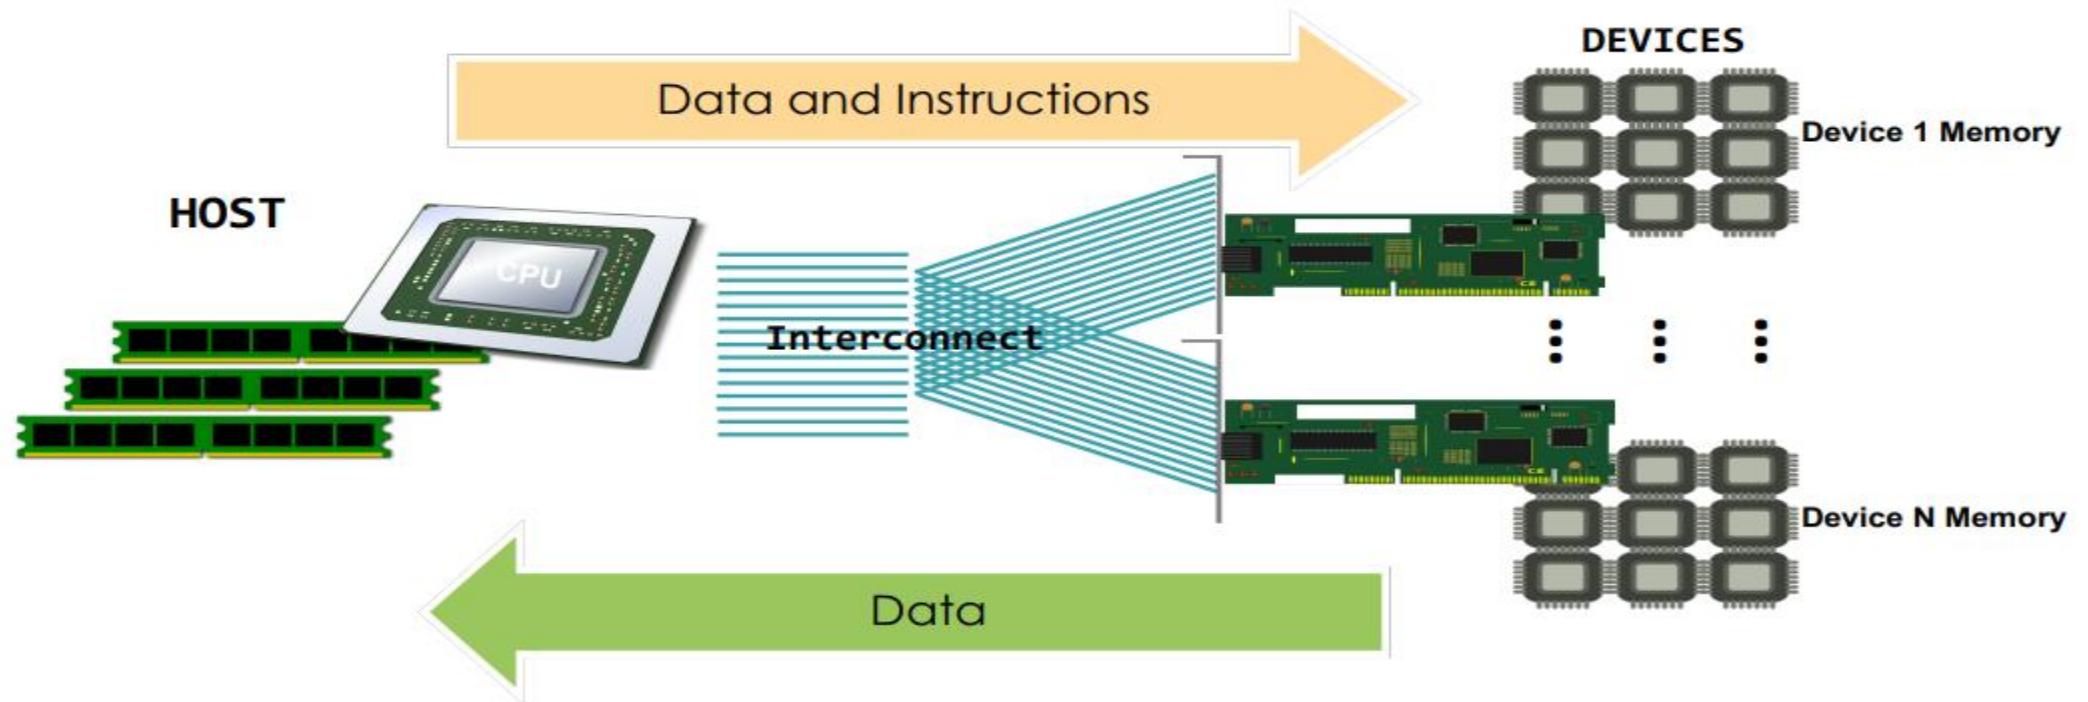
WHAT IS A GPU vs CPU?

And why GPUs are used for Machine Learning



Introduction: OpenMP Offload

OpenMP offload constructs are a set of directives for C++ and Fortran that were introduced in OpenMP 4.0 and further enhanced in later versions



OpenMP Offload: Steps

- Identification of compute kernels
 - CPU initiates kernel for execution on the device
- Expressing parallelism within the kernel
- Manage data transfer between CPU and Device
 - relevant data needs to be moved from host to device memory
 - kernel executes using device memory
 - relevant data needs to be moved from device to main memory

Step 1: Identification of Kernels to Offload

- Look for compute intensive code and that can benefit from parallel Execution
 - Use performance analysis tools to find bottlenecks
- Track independent work units with well defined data accesses
- Keep an eye on platform specs
 - GPU memory is a precious resource

How to Offload ?

| C/C++ | Description |
|--|---|
| #pragma omp target [clause[[,] clause] ...] newline structured-block | The target construct offloads the enclosed code to the accelerator. |

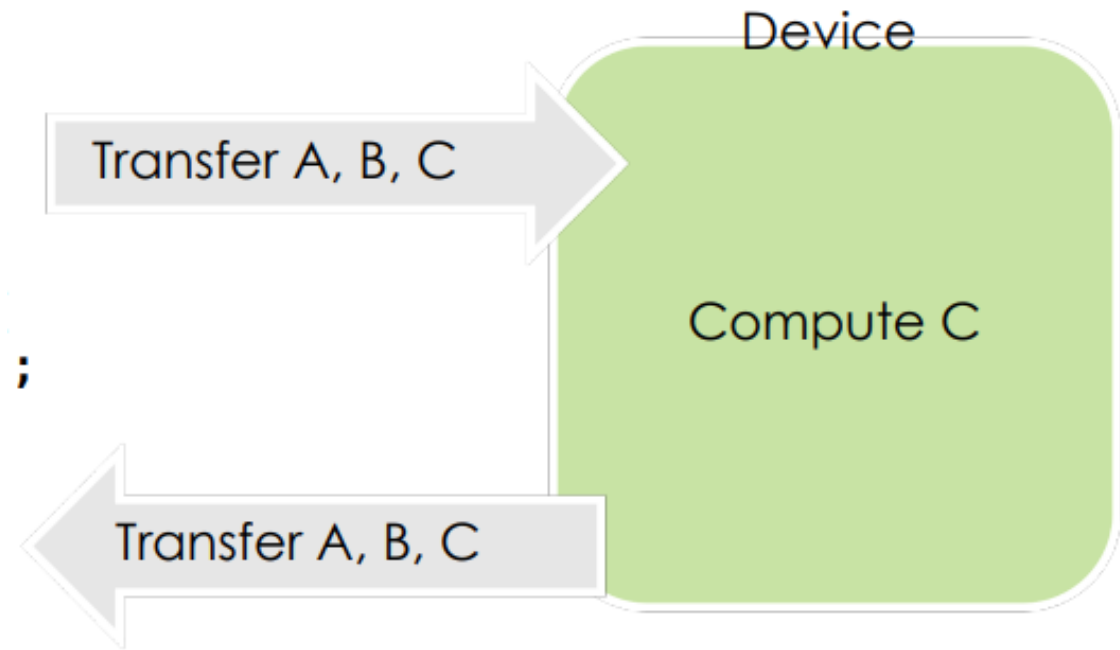
- A device data environment is created for the structured block
 - The code region is mapped to the device and executed.

Target Directive

- Clauses allowed on the target directive:
 - if([target :] scalar-expression)
 - device([device-modifier :] integer-expression)
 - thread_limit(integer-expression)
 - private(list)
 - firstprivate(list)
 - in_reduction(reduction-identifier : list)
 - map([[map-type-modifier[,] [map-type-modifier[,] ...]] map-type:] locator-list)
 - is_device_ptr(list)
 - has_device_addr(list)
 - defaultmap(implicit-behavior[:variable-category])
 - nowait
 - depend([depend-modifier,] dependence-type : locator-list)
 - allocate([allocator :] list)
 - uses_allocators(allocator[(allocator-traits-array)] [,allocator[(allocator-traits-array)] ...])

Example using omp target

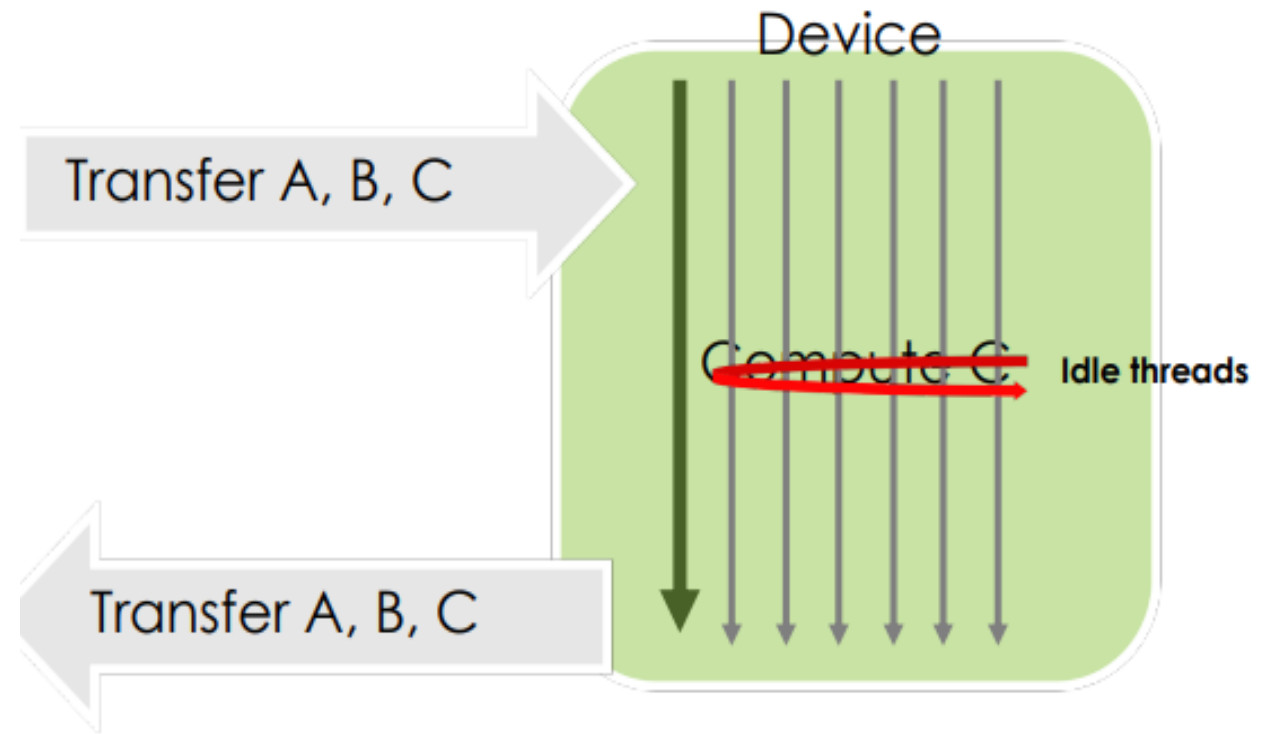
```
int A[N][N], B[N][N], C[N][N];  
/* initialize arrays */  
#pragma omp target  
{  
    for (int i = 0; i < N; ++i) {  
        for (int j = 0; j < N; ++j) {  
            C[i][j] = A[i][j] + B[i][j];  
        }  
    }  
} // end target
```



The target construct is a task generating construct

Example using omp target(Cont.)

```
int A[N][N], B[N][N], C[N][N];  
/*  
initialize arrays  
*/  
#pragma omp target  
{  
  for (int i = 0; i < N; ++i) {  
    for (int j = 0; j < N; ++j) {  
      C[i][j] = A[i][j] + B[i][j];  
    }  
  }  
} // end target
```




Device Execution Directives

| C/C++ | Description |
|---|---|
| <code>#pragma omp target [clause[[,] clause] ...] new-line structured-block</code> | The target construct offloads the enclosed code to the accelerator. |
| <code>#pragma omp target teams [clause[[,] clause] ...] new-line structured-block</code> | The target construct offloads the enclosed code to the accelerator. The teams construct creates a league of teams. The initial thread of each team executes the code region. |
| <code>#pragma omp target teams distribute [clause[[,] clause] ...] new-line loop-nest</code> | The target construct offloads the enclosed code to the accelerator. A league of thread teams is created, and loop iterations are distributed and executed by the initial teams. |
| <code>#pragma omp target teams distribute parallel for [clause[[,] clause] ...] new-line loop-nest</code> | The target construct offloads the enclosed code to the accelerator. A league of thread teams are created, and loop iterations are distributed and executed in parallel by all threads of the teams. |

Device Execution Directives(Cont.)


target

```
#pragma omp target
for (int i = 0; i < 12; ++i)
{
    C[i] = A[i] + B[i];
}
```




target teams

```
#pragma omp target teams
num_teams(3)
for (int i = 0; i < 12; ++i)
{
    C[i] = A[i] + B[i];
}
```




target teams distribute

```
#pragma omp target teams
distribute num_teams(3)
for (int i = 0; i < 12; ++i)
{
    C[i] = A[i] + B[i];
}
```



target teams distribute parallel

```
#pragma omp target teams
distribute parallel for
num_teams(3)
for (int i = 0; i < 12; ++i)
{
    C[i] = A[i] + B[i];
}
```



A large yellow arrow pointing to the right, centered on a white background. The arrow has a black outline and a solid yellow fill.

Demo



Thank You

Contact:

Name: Parikshit Ardhapurkar

Email: aparikshit@cdaca.in

Mod: 8605045054